

## 1. How do Kubernetes and Docker complement each other?

---

Docker packages apps into containers.

Kubernetes orchestrates and manages container lifecycle, networking, scaling, and updates.

Docker is the runtime used in many Kubernetes clusters (though alternatives exist).

## 2. Explain Kubernetes Architecture / Explain the Kubernetes architecture and the roles of its components.

---

**Master Node Components:**

**API Server:** Frontend to Kubernetes control plane, exposes Kubernetes API.

**etcd:** Distributed key-value store to hold cluster state data.

**Controller Manager:** Runs controllers like Node Controller, Replication Controller.

**Scheduler:** Assigns pods to nodes based on resource availability.

**Worker Node Components:**

**Kubelet:** Agent running on nodes, ensures containers are running in Pods.

**Kube-proxy:** Network proxy for service abstraction, manages network rules.

**Container Runtime:** Runs containers (Docker, containerd, CRI-O).

**Step-by-step:**

API Server processes user commands from CLI or UI. -----> Scheduler places pods on nodes -----> Controllers ensure desired state is met (like replica count). -----> Kubelet communicates with API server to manage pods on nodes.

## 3. Explain Pod Lifecycle and Management / Describe the Pod lifecycle and how Kubernetes manages Pods.

---

**Pod Lifecycle phases:** Pending → Running → Succeeded/Failed → Unknown

----- Kubernetes manages Pods -----

Kubelet creates the pod containers using container runtime. -----> If pod dies, Controllers recreate pods (if it is under ReplicaSet/Deployment). -----> Liveness and readiness probes help Kubernetes know if pod is healthy and ready.

## 4. How do you perform rolling updates and rollbacks in Kubernetes?

---

When we use `kubectl apply -f deployment.yaml` to update deployment Kubernetes creates new ReplicaSets with the new version.

The Rolling update replaces pods gradually (default `maxUnavailable=25%`, `maxSurge=25%`).

For Rollback: `kubectl rollout undo deployment/<deployment-name>`

**Example Step-by-step:**

Modify image/version in deployment manifest. -----> Apply changes using `kubectl` -----> new pods start before old pods terminate. -----> Monitor rollout status: `kubectl rollout status deployment/<name>` -----> Undo if needed to previous stable version.

+++++ Explanation ++++++

**Rolling Update:** Gradually replaces Pods of the old version with Pods of the new version without downtime.

**Rollback:** Reverts the Deployment to a previous version in case the new version has issues.

**Step 1: Create an initial Deployment** Let's create a Deployment with an NGINX container version 1.14.2.

```
# nginx-deployment.yaml

apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
          ports:
            - containerPort: 80
```

Apply this Deployment  
`kubectl apply -f nginx-deployment`  
 check the pods  
`kubectl get pod -l app=nginx`

### Step 2: Let's Perform a rolling update,

For an example Let's update the image version from `nginx:1.14.2` to `nginx:1.16.1` in the YAML file and apply the deployment or do it directly using the command:

[kubectl set image deployment/nginx-deployment nginx=nginx:1.16.1](#)

Kubernetes will start replacing pods gradually with new pods using the new image version without downtime.

Check rollout status: `kubectl rollout status deployment/nginx-deployment`

Step 3: Verify the update by checking the pods' image versions:

`kubectl get pods -l app=nginx -o jsonpath=".items[*].spec.containers[0].image"`

You should see all pods running `nginx:1.16.1`.

Step 4: Rollback to previous version if needed

If the new version has problems, rollback to previous version: `kubectl rollout undo deployment/nginx-deployment`

Check rollout status again: `kubectl rollout status deployment/nginx-deployment`

check rollout history and status or View rollout history: `kubectl rollout history deployment/nginx-deployment`

View details about a specific revision: `kubectl rollout history deployment/nginx-deployment --revision=1`

## 5. Networking in Kubernetes / Q: Explain how networking works in Kubernetes.

---

K8 Networking Aspect	What it Does	Example Resource
Pod-to-Pod communication	Direct IP communication between Pods	Pod
Stable network endpoint	Expose group of Pods via DNS/IP	Service (ClusterIP)
External access	Access Service from outside	Service (NodePort/LoadBalancer)
Network control	Restrict traffic flow between Pods	NetworkPolicy

---

External access      Uses NodePort/LoadBalancer; still may involve kube-proxy but for outside-in traffic.

Stable network endpoint      Handled via Services. Kube-proxy manages IP routing to backends

++++++

### STEP 1 -- Pod-to-Pod Direct Communication

List all pods and get their IPs: `kubectl get pods -o wide`

Example output:

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE
pod-a	1/1	Running	0	10m	10.244.1.5	node-1
pod-b	1/1	Running	0	8m	10.244.2.7	node-2

Here, pod-a has IP `10.244.1.5`, and pod-b has IP `10.244.2.7`.

Pick a pod, exec into it: `kubectl exec -it <pod-name> -- sh`

`kubectl exec -it pod-a -- sh`      You are now "inside" pod-a's container.

Try ping or curl another pod using its IP: `ping <other-pod-ip>`

`ex → ping 10.244.2.7`

You should see responses like:

`PING 10.244.2.7 (10.244.2.7): 56 data bytes`

`64 bytes from 10.244.2.7: icmp_seq=0 ttl=64 time=0.123 ms`

`64 bytes from 10.244.2.7: icmp_seq=1 ttl=64 time=0.100 ms`

Or

if it's an HTTP server running on pod-b, you can:

`ex → curl http://10.244.2.7:8080 → you will get a response directly`

You'll see that pod-to-pod communication works within the cluster without NAT.

## Step 2 → Service Discovery / Services Provide Stable Network Endpoints

---

- ☒ Pods are ephemeral and may be recreated with different IPs.
- ☒ Kubernetes Services provide a stable IP and DNS name to access a group of Pods.
- ☒ Services load balance traffic to Pods matching a label selector.

Create a Pod running nginx:

```
# nginx-pod.yaml
apiVersion: v1
kind: Pod
metadata:
  name: nginx-pod
  labels:
    app: nginx
spec:
  containers:
    - name: nginx
      image: nginx
      ports:
        - containerPort: 80
```

STEP - 1

Create a Service to expose nginx Pod inside the cluster:

```
# nginx-service.yaml
apiVersion: v1
kind: Service
metadata:
  name: nginx-service
spec:
  selector:
    app: nginx
  ports:
    - protocol: TCP
      port: 80      # Service port
      targetPort: 80 # Pod container port
```

STEP - 2

Apply both manifest : `kubectl apply -f nginx-pod.yaml` and `kubectl apply -f nginx-service.yaml`

Check Service IP: `kubectl get svc nginx-service`

Output:

NAME	TYPE	CLUSTER-IP	PORT(S)	AGE
nginx-service	ClusterIP	10.96.25.100	80/TCP	1m

You can now access nginx via this stable IP or DNS name `nginx-service` from other Pods inside the cluster.

## Accessing Services from other pod Inside the Cluster

---

Lets Create a Pod with a curl client to test access:

```
# curl-pod.yaml
apiVersion: v1
kind: Pod
metadata:
  name: curl-pod
spec:
  containers:
    - name: curl
      image: curlimages/curl
      command: ["sleep", "3600"]
```

Apply: `kubectl apply -f curl-pod.yaml`

Exec into curl-pod and curl nginx service by DNS name:

`kubectl exec -it curl-pod -- curl http://nginx-service`

You should get the nginx welcome page HTML output.

### Step 3 → Exposing Services Outside the Cluster

---

To allow access from outside, change Service type to `NodePort` or `LoadBalancer`.

Example `NodePort` service:

```
apiVersion: v1
kind: Service
metadata:
  name: nginx-nodeport
spec:
  selector:
    app: nginx
  ports:
    - port: 80
      targetPort: 80
      nodePort: 30080
  type: NodePort
```

```
kubectl apply -f nginx-nodeport.yaml
```

```
kubectl get nodes -o wide
```

Now you can access nginx from outside by hitting:

```
http://<Node-IP>:30080
```

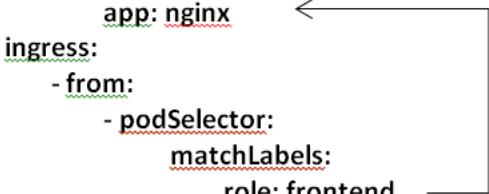
### Step 4 → Controlling Traffic with Network Policies

---

By default, all Pods can talk to each other. You can create `NetworkPolicies` to restrict which Pods can talk to which.

Example allowing only Pods labeled `role=frontend` to access nginx on port 80:

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: allow-frontend
spec:
  podSelector:
    matchLabels:
      app: nginx
  ingress:
    - from:
        - podSelector:
            matchLabels:
              role: frontend
  ports:
    - protocol: TCP
      port: 80
```



### 6. Explain the difference between Deployment, ReplicaSet, StatefulSet, and DaemonSet Controllers.

---

Answer:

**Deployment:** Manages stateless apps, handles rolling updates and rollbacks.

**ReplicaSet:** Ensures specified number of pod replicas are running, usually managed by Deployments.

**StatefulSet:** Manages stateful apps, provides stable identities and persistent storage.

**DaemonSet:** Ensures a copy of a pod runs on all (or some) nodes, e.g., for logging or monitoring agents.

### 7. Explain how Kubernetes handles storage?

---

Kubernetes supports Volumes which persist data beyond the lifecycle of a Pod. There are different types of volumes: `emptyDir`, `hostPath`, `persistentVolumeClaim (PVC)`, etc.

**PersistentVolumes (PV)** represent actual storage resources. It is a piece of storage in the cluster provisioned by an admin or dynamically provisioned using a `StorageClass` (e.g., SSD vs HDD, reclaim policy, etc.).

**PersistentVolumeClaim (PVC):** A user's request for storage.

### 1. Define a StorageClass

```
# storage-class.yaml
apiVersion: storage.k8s.io/v1
kind: StorageClass
metadata:
  name: standard
provisioner: kubernetes.io/aws-ebs # Use a provisioner
parameters: (e.g., aws-ebs, gce-pd, csi)
  type: gp2
reclaimPolicy: Retain
volumeBindingMode: WaitForFirstConsumer
```

### 2. Create a PersistentVolumeClaim (PVC)

```
# pvc.yaml
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: my-pvc
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 1Gi
  storageClassName: standard
```

Apply it: `kubectl apply -f storage-class.yaml` and `kubectl apply -f pvc.yaml`

This triggers dynamic provisioning using the StorageClass. Kubernetes creates a PersistentVolume behind the scenes.

Now Create a Pod which will use the PVC

```
# pod_using_pvc.yaml
apiVersion: v1
kind: Pod
metadata:
  name: pvc-demo-pod
spec:
  containers:
    - name: app
      image: busybox
      command: ["sleep", "3600"]
      volumeMounts:
        - mountPath: "/data"
          name: storage
  volumes:
    - name: storage
      persistentVolumeClaim:
        claimName: my-pvc
```

Apply it: `kubectl apply -f pod_using_pvc.yaml`

Now your container can write to `/data`, and data will persist as long as the volume exists, even if the Pod is deleted.

```
kubectl get pvc
kubectl get pv
kubectl get pod pvc-demo-pod -o yaml
```

## 8. What are Controllers in Kubernetes?

---

Controllers monitor the state of the cluster via the API server and make or request changes to move the current state towards the desired state.

They automate tasks such as: 1) Ensuring a certain number of pods are running , 2) Restarting failed containers , 3) Managing updates , 4) Scaling applications

### Types of Built-in Controllers

ReplicaSet Controller –	Ensures a specific number of Pod replicas are running.
Job Controller –	Ensures Pods complete a specific task to completion.
DaemonSet Controller –	Ensures a copy of a Pod runs on all or some Nodes.
StatefulSet Controller –	Manages stateful applications (with persistent identity).
CronJob Controller –	Manages time-based jobs.

### Step 1: Define the ReplicaSet Controller

---

This YAML defines a Deployment that manages 3 replicas of an NGINX web server.

```
# nginx-deployment.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: my-nginx
          image: nginx:1.25
          ports:
            - containerPort: 80
```

**Apply the Deployment YAML : `kubectl apply -f nginx-deployment.yaml`**

**Check the Deployment Status `kubectl get deployments`**

You'll see something like:

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
nginx-deployment	3/3	3	3	10s

-----> This means  
the Deployment controller created 3 Pods via a ReplicaSet.

**Verify the Pods**

```
kubectl get pods -l app=nginx
Should return 3 pods like:
nginx-deployment-xxxxx  Running
nginx-deployment-yyyyy  Running
nginx-deployment-zzzzz  Running
```

## Step 2 → Define StatefulSet Controller

**StatefulSet Controller** manages stateful applications — those needing stable identities, persistent storage, and ordered deployment. Each pod gets a unique name and persistent volume. Useful for databases (like MySQL, Cassandra, etc.)

```
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: web
spec:
  serviceName: "nginx"
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx
          volumeMounts:
            - name: www
              mountPath: /usr/share/nginx/html
  volumeClaimTemplates:
    - metadata:
        name: www
    spec:
      accessModes: [ "ReadWriteOnce" ]
      resources:
        requests:
          storage: 1Gi
```

## Step 5 → CronJob Controller

Runs Jobs on a schedule, like a Linux cron job.  
Can manage retries and failed runs. Syntax follows standard cron format.

```
apiVersion: batch/v1
kind: CronJob
metadata:
  name: hello-cron
spec:
  schedule: "*/1 * * * *"          # Every minute
  jobTemplate:
    spec:
      template:
        spec:
          containers:
            - name: hello
              image: busybox
              command: ["echo", "Hello from CronJob!"]
            restartPolicy: OnFailure
```

Use-case: Periodic backups, reporting, monitoring jobs.

Use-case: Databases, message queues, or anything needing stable storage and identity.

### Step 3 → DaemonSet Controller

Ensures that a copy of a Pod runs on every / selected Node in the cluster. Commonly use for logging, monitoring, or network agents. Use-case: Run agents like Fluentd, Prometheus Node Exporter, etc.

```
apiVersion: apps/v1
kind: DaemonSet
metadata:
  name: node-exporter
spec:
  selector:
    matchLabels:
      name: node-exporter
  template:
    metadata:
      labels:
        name: node-exporter
    spec:
      containers:
        - name: node-exporter
          image: prom/node-exporter
```

Use-case: Run agents like Fluentd, Prometheus Node Exporter, etc.

### Step 4 → Job Controller

Creates Pods that run to completion (once, not forever).  
Perfect for batch jobs or one-time tasks. Retries if Pods fail.

```
apiVersion: batch/v1
kind: Job
metadata:
  name: hello-job
spec:
  template:
    spec:
      containers:
        - name: hello
          image: busybox
          command: ["echo", "Hello from Job!"]
          restartPolicy: Never
          backoffLimit: 4 -----> 4 retries if pod fails
```

Use-case: One-off tasks like database migrations, data processing

## 9. How do you secure a Kubernetes cluster?

Use RBAC (Role-Based Access Control) for permission management.

Enable network policies to control pod-to-pod communication.

Use TLS for communication between components.

Use image scanning and trusted registries.

Enable Pod Security Policies or Pod Security Admission for runtime restrictions.

Encrypt secrets at rest.

Regularly patch and update cluster components.

### Step 1. Use Role-Based Access Control (RBAC)

RBAC use to limit who can perform which action in your cluster. For that we need to create RBAC role and bind it to the user.

Let's create a role that allows reading pods:

```
apiVersion: rbac.authorization.k8s.io/v1
kind: Role
metadata:
  namespace: default
  name: pod-reader
rules:
  - apiGroups: []
    resources: ["pods", "services"]
    verbs: ["get", "watch", "list"]
```

1<sup>st</sup> step

Now bind role to a user:

```
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: read-pods-binding
  namespace: default
subjects:
  - kind: User
    name: dev-group
    apiGroup: rbac.authorization.k8s.io
roleRef:
  kind: Role
  name: pod-reader
  apiGroup: rbac.authorization.k8s.io
```

2<sup>nd</sup> step

Apply: kubectl apply -f role.yaml and kubectl apply -f rolebinding.yaml

**Enable Authentication and Authorization using client certificates, OIDC, or cloud provider IAM.**

❑ Use OIDC with Google or Azure AD to authenticate users.

❑ Configure --authentication-mode=Webhook and --authorization-mode=RBAC in the API server.  
Step 2 → Use Network Policies

---

Control traffic between pods and namespaces, by enabling a network plugin that supports policies (e.g., Calico, Cilium)

`kubectl apply -f https://docs.projectcalico.org/manifests/calico.yaml` → Download and apply Calico manifests  
`kubectl get pods -n calico-system` → Verify Calico pods are running

Block all traffic except from a specific app using NetworkPolicy :

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: allow-nginx
  namespace: default
spec:
  podSelector:
    matchLabels:
      app: nginx
  policyTypes:
    - Ingress
  ingress:
    - from:
        - podSelector:
            matchLabels:
              app: frontend
```

Step 3 → Use TLS for communication between components.  
As etcd contains all cluster secrets and configurations.

- ❑ Encrypt etcd at rest.
- ❑ Use TLS for etcd client and peer connections.
- ❑ Limit access to etcd to the API server only.

Vi `/etc/kubernetes/encryption-config.yaml`

```
apiVersion: apiserver.config.k8s.io/v1
kind: EncryptionConfiguration
resources:
  - resources:
      - secrets
  providers:
    - aescbc:
        keys:
          - name: key1
            secret: <base64-encoded-secret>
    - identity: {}
```

Edit `/etc/kubernetes/manifests/kube-apiserver.yaml` file and  
add the --encryption-provider-config=`/etc/kubernetes/encryption-config.yaml`  
Restart the `kube-apiserver` if needed

**Step 4 → Scan Container Images**

Use tools like Trivy, Clair, or Aqua Security scan the container image

Ex → `trivy image nginx:latest`

## 10. What strategies are available for deploying applications in Kubernetes?

---

**Rolling Update:** Incrementally updates pods with zero downtime.

**Blue/Green Deployment:** Runs two environments and switches traffic from old to new.

**Canary Deployment:** Gradually rolls out new versions to a subset of users.

**Recreate:** Shuts down old version before starting the new one.

**Rolling update** → explained at above

**Blue/green deployment** →

In case of blue/green deployment the downtime and risk are reduced by running two production environments (Blue and Green) in parallel. Once test is success we can delete old version.

Lets think the Current version is running in the Blue environment. Green is a new version of your app, deployed alongside Blue. Then tests/health checks are run against Green. Once health check is good switch traffic from Blue to Green via Ingress or LoadBalancer. Once confirmed, Blue is optionally deleted or kept for rollback.

To perform blue/green deployment there should be two deployment manifest & a selector base service manifest file.

Vi deployment-blue.yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: myapp-blue
spec:
  replicas: 3
  selector:
    matchLabels:
      app: myapp
      version: blue
  template:
    metadata:
      labels:
        app: myapp
        version: blue
    spec:
      containers:
        - name: myapp
          image: myapp:1.0
      ports:
        - containerPort: 80
```

Vi deployment-green.yaml

Same as above, but the version: green and image: new\_image\_with\_new\_version as below.

```
metadata:
  name: myapp-green
...
version: green
image: myapp:2.0
```

vi selector\_base\_service.yaml

```
apiVersion: v1
kind: Service
metadata:
  name: myapp-service
spec:
  selector:
    app: myapp
    version: blue    # Change to 'green' during rollout
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
```

vi blue\_green\_deploy.yml ----- it is in Ansible

```
- name: Blue-Green Deployment on Kubernetes
hosts: localhost
gather_facts: no
vars:
  new_version: green # or blue
  old_version: blue # or green
  image_tag: "2.0"

tasks:
  - name: Apply new version deployment
    kubernetes.core.k8s:
      state: present
      definition: "{{ lookup('file', 'deployment-' + new_version + '.yaml') }}"

  - name: Wait for rollout to complete
    command: >
      kubectl rollout status deployment/myapp-{{ new_version }}
    register: rollout_status
    until: rollout_status.stdout.find('successfully rolled out') != -1
    retries: 10
    delay: 10
```

```

- name: Run health checks (optional)
  uri:
    url: "http://{{ new_version }}.myapp.internal/health"
    status_code: 200
  register: health_result
  until: health_result.status == 200
  retries: 5
  delay: 5
- name: Switch service to new version
  kubernetes.core.k8s:
    state: present
    definition:
      apiVersion: v1
      kind: Service
      metadata:
        name: myapp-service
      spec:
        selector:
          app: myapp
          version: "{{ new_version }}"
        ports:
          - protocol: TCP
            port: 80
            targetPort: 80

- name: (Optional) Delete old version deployment
  kubernetes.core.k8s:
    state: absent
    kind: Deployment
    name: myapp-{{ old_version }}

```

**Canary Deployment → Canary deployment** is a process to roll out a new version of an application to a small subset of users before exposing it to the entire user base. This lets you test the new version in production with real traffic, minimizing risk by limiting exposure to potential issues.

For Canary deployment we need to follow below steps

Existing app version (stable).

New app version (canary) deployed alongside stable.

Traffic is split between stable and canary pods, e.g., 90% stable, 10% canary.

You monitor canary behavior.

If successful, shift 100% traffic to new version and remove old version.

```

ansi-canary-folder/
├── canary-deployment.yaml
├── stable-deployment.yaml
├── service.yaml
└── canary-playbook.yaml

```

```
stable-deployment.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: myapp-stable
spec:
  replicas: 9
  selector:
    matchLabels:
      app: myapp
      version: stable
  template:
    metadata:
      labels:
        app: myapp
        version: stable
    spec:
      containers:
        - name: myapp
          image: mydockerhub/myapp:v1
          ports:
            - containerPort: 80
```

```
canary-deployment.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: myapp-canary
spec:
  replicas: 1
  selector:
    matchLabels:
      app: myapp
      version: canary
  template:
    metadata:
      labels:
        app: myapp
        version: canary
    spec:
      containers:
        - name: myapp
          image: mydockerhub/myapp:v2
          ports:
            - containerPort: 80
```

```
canary_service.yaml
apiVersion: v1
kind: Service
metadata:
  name: myapp-service
spec:
  selector:
    app: myapp
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
  type: LoadBalancer
```

#### Ansible Playbook - canary-playbook.yaml

```
- name: Canary Deployment on Kubernetes EKS
hosts: localhost
gather_facts: no
tasks:

  - name: Deploy Stable version
    k8s:
      state: present
      definition: "{{ lookup('file', 'stable-deployment.yaml') }}"

  - name: Deploy Canary version
    k8s:
      state: present
      definition: "{{ lookup('file', 'canary-deployment.yaml') }}"

  - name: Deploy Service exposing both versions
    k8s:
      state: present
      definition: "{{ lookup('file', 'service.yaml') }}"

  - name: Wait for stable deployment rollout to complete
    k8s_info:
      kind: Deployment
      namespace: default
      name: myapp-stable
    register: stable_deploy

  - name: Wait until stable replicas ready
    k8s_info:
      kind: Pod
      namespace: default
      label_selectors:
```

```

- "app=myapp,version=stable"
register: stable_pods
until: stable_pods.resources | selectattr('status.phase','equalto','Running') | list | length == 9
retries: 10
delay: 15

- name: Wait until canary pods ready
k8s_info:
  kind: Pod
  namespace: default
  label_selectors:
    - "app=myapp,version=canary"
register: canary_pods
until: canary_pods.resources | selectattr('status.phase','equalto','Running') | list | length == 1
retries: 10
delay: 15

```

Once real time testing complete by users then scaling Canary Up/Stable Down (Promoting Canary) using Ansible playbook

```

- name: Promote canary to stable by scaling
hosts: localhost
gather_facts: no

tasks:
- name: Scale stable deployment down
k8s:
  kind: Deployment
  name: myapp-stable
  namespace: default
  replicas: 0

- name: Scale canary deployment up
k8s:
  kind: Deployment
  name: myapp-canary
  namespace: default
  replicas: 10

- name: Update canary label to stable (optional)
k8s:
  kind: Deployment
  name: myapp-canary
  namespace: default
  definition:
    spec:
      template:
        metadata:
          labels:
            version: stable

```

## 11. How do you monitor a Kubernetes cluster?

---

Use tools like Prometheus (metrics collection), Grafana (visualization).

Use ELK/EFK stacks (Elasticsearch, Fluentd/Fluent Bit, Kibana) for logs.

Kubernetes Dashboard or Lens IDE for cluster state.

Use probes (liveness/readiness) for pod health.

## 12. Explain how you would troubleshoot a failing pod?

---

Check pod status:

`kubectl get pods`

Check error message and exit status of pod in pod logs:

`kubectl logs <pod-name>`

`kubectl logs myapp-pod -c <container-name>` -----> If multiple containers

Inspect events for failures:

`kubectl describe pod <pod-name>`

`kubectl get events --sort-by='lastTimestamp'`

Check container runtime status on the node ( MemoryPressure and DiskPressure will be True)

`kubectl describe node nodeName`

Validate resource limits/requests and node capacity.

`kubectl get pod <pod-name> -o yaml`

ex→

resources:

  requests:

    cpu: "4"

    memory: "8Gi"

If your node only has 2 CPUs, this pod will remain in Pending state.

Check network connectivity and service endpoints.

`Kubectl get svc`

`Kubectl get endpoints` -----> it will give ip and port use that port in below command

`Kubectl exec -it <pod-name> -- curl <service>:<port>`

## 13. What is a ConfigMap and Secret? How do you use them?

---

**ConfigMap:** Used to store non-sensitive configuration data as key-value pairs.

**Secret:** Used to store sensitive data like passwords, tokens, or keys, base64-encoded.

They can be consumed by Pods as environment variables or mounted as files.

First create the encrypted data

Create configmap from yaml file

```
echo -n 'admin' | base64 -----> o/p → YWRtaW4
echo -n 'suman@#345' | base64 -----> o/p → MWETY345DAV45AK
```

Vi testConfigMap.yaml

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: my-config
data:
  APP_ENV: "production"
  APP_DEBUG: "false"
  DATABASE_URL: "mysql://db:3306"
```

Create Secret from yaml file

```
Vi mySecretFile.yaml
apiVersion: v1
kind: Secret
metadata:
  name: my-secret
type: Opaque
data:
  DB_PASSWORD: MWETY345DAV45AK
```

Create configmap

`kubectl apply -f testConfigMap.yaml`

Create secret

`kubectl apply -f mySecretFile.yaml`

```
Use this ConfigMap in a Pod as environment variables
apiVersion: v1
kind: Pod
metadata:
  name: configmap-demo
spec:
  containers:
    - name: app
      image: myapp:latest
      envFrom:
        - configMapRef:
            name: my-config
```

Create a property file  
`cat /etc/config/app.properties`

```
APP_ENV=production
APP_DEBUG=false
DATABASE_URL=mysql://db:3306
```

Create a Configmap yaml  
`Vi my_configmap.yaml`

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: app-config
data:
  app.properties:
    APP_ENV=production
    APP_DEBUG=false
    DATABASE_URL=mysql://db:3306
```

Or Create a Configmap from the command

`Kubectl create configmap app-config --from-file=/etc/config/app.properties`

```
Use this Secret in a Pod as environment variables
apiVersion: v1
kind: Pod
metadata:
  name: secret-demo
spec:
  containers:
    - name: app
      image: myapp:latest
      env:
        - name: DB_PASSWORD
          valueFrom:
            secretKeyRef:
              name: my-secret
              key: DB_PASSWORD
```

Define the Pod that mounts the ConfigMap as a volume

```
apiVersion: v1
kind: Pod
metadata:
  name: configmap-volume-pod
spec:
  containers:
    - name: app-container
      image: busybox
      command: [ "sleep", "3600" ]
      volumeMounts:
        - name: config-volume
          mountPath: /etc/config # Files from ConfigMap(app.properties) is available here
  volumes:
    - name: config-volume
      configMap:
        name: app-config # Refers to the ConfigMap created earlier
```

Step-1

Step-2

First create the encrypted data

```
echo -n 'admin' | base64 -----> o/p → YWRtaW4
echo -n 'suman@#345' | base64 -----> o/p → MWETY345DAV45AK
```

Now create the manifest file

`Vi mySecret.yaml`

Step 1

```
apiVersion: v1
kind: Secret
metadata:
  name: mysecret
  type: Opaque
data:
  username: YWRtaW4 → Here username is the key and encrypted data is value
  password: MWETY345DAV48AK → password is key and encrypted data is value
```

create pod manifest file to use secret

```
apiVersion: v1
kind: Pod
metadata:
  name: secret-volume-pod
spec:
  containers:
    - name: secret-container
      image: busybox
      command: [ "sleep", "3600" ]
      volumeMounts:
        - name: secret-volume
          mountPath: /etc/secret-data
          readOnly: true
  volumes:
    - name: secret-volume
      secret:
        secretName: my-secret
```

Step 2

`Kubectl create -f mySecretPod.yaml`

`Kubectl exec secret-volume-pod ls /etc/secret-data` → two files username , password will be there

## 14. What is Helm and why would you use it?

---

Helm is a package manager for Kubernetes, simplifying app deployment and management.  
It uses Charts to define, install, and upgrade Kubernetes applications.  
It helps manage complex apps and their dependencies declaratively.

## 15. Deploy NGINX App with Helm?

---

Step 1 → Create the folder structure / Skeleton for Helm Chart

`helm create my-nginx-test` → it will create below folder structure

```
my-nginx-test/
  Chart.yaml      → Defines metadata about the chart (name, version, chart dependencies)
  values.yaml     → it is used to pass label to yaml
  templates/
    deployment.yaml
    service.yaml
    _helpers.tpl
```

Step 2 → create `values.yaml`

```
replicaCount: 2
image:
  repository: nginx
  tag: latest
  pullPolicy: IfNotPresent
service:
  type: ClusterIP
  port: 80
resources: {}
nodeSelector: {}
tolerations: []
affinity: {}
```

```
my-nginx/          # Github repo
  .github/
  helm/
    my-nginx-test/      # Helm chart folder
    |   Chart.yaml       # Chart metadata
    |   values.yaml      # Default configuration values
    |   templates/        # Kubernetes YAML templates
    |     deployment.yaml
    |     service.yaml
    |     _helpers.tpl
    |   README.md        # Chart-specific README
    nginx.conf          # NGINX configuration file
    Dockerfile          # To build custom NGINX image (optional)
    README.md           # Project overview and instructions
    LICENSE             # License file (e.g., MIT)
    .dockerignore        # Ignore files for Docker build
```

Step 3 → `_helpers.tpl` is the file used to store reusable template helpers

```
{{/* Generate chart name */}}
{{- define "my-nginx-test.name" -}}
{{- .Chart.Name -}}
{{- end }}

{{/* Generate full resource name */}}
{{- define "my-nginx-test.fullname" -}}
{{- printf "%s-%s" .Release.Name .Chart.Name | trunc 63 | trimSuffix "-" -}}
{{- end }}

{{/* Common labels for resources */}}
{{- define "my-nginx-test.labels" -}}
app.kubernetes.io/name: {{ include "my-nginx-test.name" . }}
app.kubernetes.io/instance: {{ .Release.Name }}
app.kubernetes.io/version: {{ .Chart.AppVersion }}
helm.sh/chart: {{ .Chart.Name }}-{{ .Chart.Version }}
{{- end }}
```

Step 4 → `deployment.yaml` which manages pods, specifying things like replicas, container images, ports, and labels

```

apiVersion: apps/v1
kind: Deployment
metadata:
  name: {{ include "my-nginx-test.fullname" . }}
  labels:
    {{ include "my-nginx-test.labels" . | indent 4}}
spec:
  replicas: {{ .Values.replicaCount }}
  selector:
    matchLabels:
      app.kubernetes.io/name:{{ include "my-nginx-test.name" . }}
      app.kubernetes.io/instance:{{ .Release.Name }}
  template:
    metadata:
      labels:
        app.kubernetes.io/name:{{ include "my-nginx-test.name" . }}
        app.kubernetes.io/instance:{{ .Release.Name }}
    spec:
      containers:
        - name: nginx
          image: "{{ .Values.image.repository }}:{{ .Values.image.tag }}"
          imagePullPolicy: {{ .Values.image.pullPolicy }}
        ports:
          - containerPort: {{ .Values.service.port }}
      livenessProbe:
        httpGet:
          path: /
          port: {{ .Values.service.port }}
      readinessProbe:
        httpGet:
          path: /
          port: {{ .Values.service.port }}
      resources:
        {{- toYaml .Values.resources | nindent 12 .}}

```

Step 5 → `service.yaml` exposes your pods to other services or the outside world.

```

apiVersion: v1
kind: Service
metadata:
  name: {{ include "my-nginx-test.fullname" . }}
  labels:
    app: {{ include "my-nginx-test.name" . }}
spec:
  type: {{ .Values.service.type }}
  ports:
    - port: {{ .Values.service.port }}
      targetPort: {{ .Values.service.port }}
      protocol: TCP
      name: http
  selector:
    app.kubernetes.io/name:{{ include "my-nginx-test.name" . }}
    app.kubernetes.io/instance:{{ .Release.Name }}

```

Step 6 → create `Chart.yaml`

```

apiVersion: v2          # Helm chart API version (v2 is for Helm 3+)
name: my-nginx-test    # Name of the chart
description: A simple NGINX web server
type: application       # Can be "application" or "library"
version: 1.0.0           # Chart version (used for Helm chart packages)
appVersion: "1.21.6"    # Version of app being deployed (e.g. NGINX version)
keywords:
  - nginx
  - web
  - http
maintainers:
  - name: sumanta
    email: sk@example.com
home: https://nginx.org
sources:
  - https://github.com/example/my-nginx

```

**Step 7 → Install Your Chart**

```
helm install nginx-app ./my-nginx-test-proj
```

**Step 8 → update replicaCount value to 3 in values.yaml we can upgrade using helm**

```
helm upgrade nginx-app ./ my-nginx-test-proj
```

## 16. Explain Horizontal Pod Autoscaling (HPA).

===== HPA

automatically scales the number of pods based on CPU/memory usage or custom metrics.

It queries the Metrics Server for resource usage.

Configuration involves defining min/max pod count and target utilization.

Example

### Step 1: Deploy NGINX using deployment manifest file

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  replicas: 1
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx
          resources:
            requests:
              cpu: 100m
            limits:
              cpu: 500m
```

### Step 2 : create the HPA (Horizontal Pod Autoscaling) manifest file

```
apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler
metadata:
  name: nginx-hpa
spec:
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: nginx-deployment <----- it look for deployment manifest file
  minReplicas: 1
  maxReplicas: 5
  metrics:
    - type: Resource
      resource:
        name: cpu
        target:
          type: Utilization
          averageUtilization: 50
```

kubectl apply -f nginx-deployment.yaml

kubectl apply -f nginx-hpa.yaml

kubectl get hpa nginx-hpa → Check HPA status

Generate load on NGINX to increase CPU usage (for example using kubectl exec to run a stress command inside a pod or an external load generator). The HPA will increase pod count automatically.